

See second page for diagram.

*note: I tracked the current “time” (tick number) as a field of my Scheduler class. Thus, I was able to use a hashmap with the completion time of each job as keys and the jobs themselves as values to store the running jobs. Thus, at each tick I was able to remove each completed job in constant time per job.

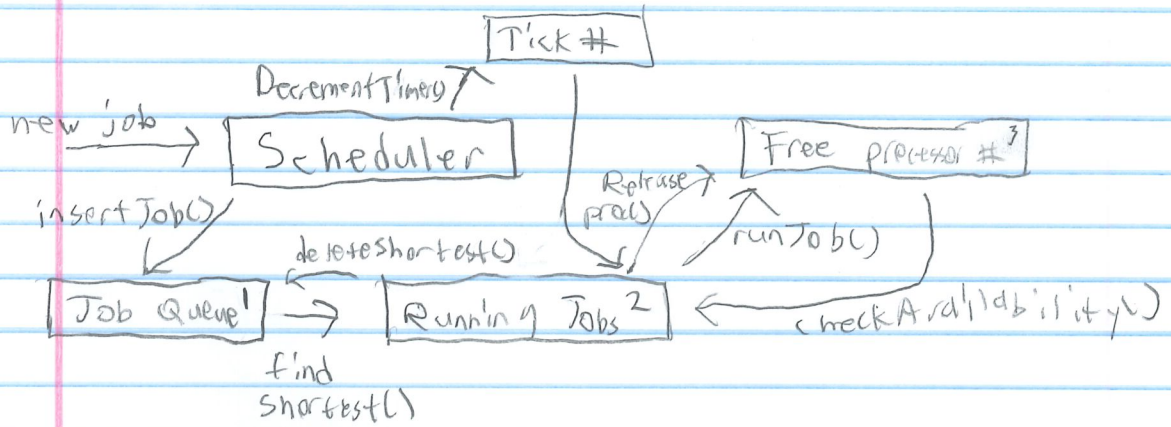
Worst-case complexities:

n is the total number of jobs.

m is the maximum number of jobs completing at the same tick.

insertJob()	$O(\log n)$
findShortest()	$O(1)$
deleteShortest()	$O(\log n)$
checkAvailability()	$O(1)$
runJob()	$O(1)$
decrementTimer()	$O(1)$
releaseProcs()	$O(m)$

The main shortcoming of this shortest-job-first strategy is that it completely ignores how many processors the job takes. For example, if there are only 5 processors available, and job A takes 5 processors and 10 ticks, and job B takes 6 processors and 9 ticks, our strategy will wait until at least 1 more processor becomes available in order to be able to run job B. Job B becomes the bottleneck, causing our strategy to waste 5 processors for an indefinite amount of ticks. If no processors would become available for 10 ticks, an optimal strategy would have completed job A before our strategy even starts job B or job A. Clearly our strategy is suboptimal in terms of both performance and functionality.



- ¹ std::make_heap()
- ² std::unordered_map
- ³ integer number